

EDAN95

Applied Machine Learning

<http://cs.lth.se/edan95/>

Lecture 9: Autoencoders and Generative Learning

Pierre Nugues


Lund University

Pierre.Nugues@cs.lth.sehttp://cs.lth.se/pierre_nugues/

December 3, 2018

The Functional Model

So far, we have used the Sequential model to build networks
These models correspond to pipelines with one input and one output



To build graphs, we need to use the functional model.

Comparing the Models

Sequential:

```
seq_model = Sequential()  
seq_model.add(layers.Dense(32, activation='relu',  
    input_shape=(64,)))  
seq_model.add(layers.Dense(32, activation='relu'))  
seq_model.add(layers.Dense(10, activation='softmax'))
```

Functional:

```
input_tensor = Input(shape=(64,))  
x = layers.Dense(32, activation='relu')(input_tensor)  
x = layers.Dense(32, activation='relu')(x)  
output_tensor = layers.Dense(10, activation='softmax')(x)  
model = Model(input_tensor, output_tensor)
```

From Chollet, page 237

Building a Multi Input Model

To build a multi input, we need the functional model and at a certain point, merge the branches with `layers.concatenate()` function
We will now build a NER tagger that uses two inputs: the words and parts of speech

The Word Branch

```
text_vocabulary_size = len(word_set) + 2
text_input = Input(shape=(None,), dtype='int32', name='text')
embedded_text = layers.Embedding(text_vocabulary_size,
                                  64, mask_zero=True)(text_input)
encoded_text = layers.LSTM(32,
                           return_sequences=True)(embedded_text)
```

The Part-of-Speech Branch

```
pos_vocabulary_size = len(pos_set) + 2
pos_input = Input(shape=(None,),
                   dtype='int32',
                   name='pos')
embedded_pos = layers.Embedding(pos_vocabulary_size,
                                32, mask_zero=True)(pos_input)
encoded_pos = layers.LSTM(16,
                          return_sequences=True)(embedded_pos)
```

Merging and Common Part

```
concatenated = layers.concatenate(  
    [encoded_text, encoded_pos], axis=-1)  
ner_vocabulary_size = len(ner_set) + 2  
ner = layers.Dense(ner_vocabulary_size,  
    activation='softmax')(concatenated)  
  
model = Model([text_input, pos_input], ner)  
model.compile(optimizer='rmsprop',  
    loss='categorical_crossentropy',  
    metrics=['acc'])
```


Code Example

The NER tagger with two inputs: the words and parts of speech and we will compare it to a sequential model

Jupyter Notebooks: `5.2-multiinput.ipynb` and `5.3-monoinput.ipynb`

Multiple Outputs

It is also possible to build a model with multiple outputs, for instance the word as input to predict the parts of speech and the named entities.



The Word Input

```
text_vocabulary_size = len(word_set) + 2
text_input = Input(shape=(None,), dtype='int32', name='text')
embedded_text = layers.Embedding(text_vocabulary_size,
                                  64, mask_zero=True)(text_input)
encoded_text = layers.LSTM(32,
                           return_sequences=True)(embedded_text)
```

The POS output

```
pos_vocabulary_size = len(pos_set) + 2
pos_output = layers.Dense(pos_vocabulary_size,
                           activation='softmax',
                           name='pos')(encoded_text)
```

The NER Output

```
ner_vocabulary_size = len(ner_set) + 2
ner_output = layers.Dense(ner_vocabulary_size,
                           activation='softmax',
                           name='ner')(encoded_text)
```

The Model



```
model = Model(text_input, [pos_output, ner_output])
model.compile(optimizer='rmsprop',
              loss=['categorical_crossentropy',
                   'categorical_crossentropy'],
              metrics=['acc'])

model.fit(X_words_idx,
        {'pos':Y_pos_idx_cat, 'ner':Y_ner_idx_cat},
        epochs=3, batch_size=128)
```

It is possible to build more complex models, provided that they have the form of a directed acyclic graph. See the book.

Monitoring Training

We have seen different shapes of validation accuracies and loss:



15 epochs seem the optimal number and it is probably useless to run more. Keras provided callbacks for this.

Two Callbacks

- 1 `keras.callbacks.EarlyStopping` to stop training when validation scores do not improve;
- 2 `keras.callbacks.ModelCheckpoint` to save models

```
callbacks_list = [  
    keras.callbacks.EarlyStopping(  
        monitor='acc',  
        patience=1,),  
    keras.callbacks.ModelCheckpoint(  
        filepath='my_model.h5',  
        monitor='val_loss',  
        save_best_only=True,)  
]
```

From Chollet, page 250

Including the Callbacks

```
model.compile(optimizer='rmsprop',  
              loss='binary_crossentropy',  
              metrics=['acc'])  
  
model.fit(x, y,  
          epochs=10,  
          batch_size=32,  
          callbacks=callbacks_list,  
          validation_data=(x_val, y_val))
```

You can also write your own callbacks, see Chollet, page 251-252

Tensorboard

Tensorboard is a visualization tool

You include it with a callback

```
callbacks = [  
    keras.callbacks.TensorBoard(  
        log_dir='tb_log_folder',  
        histogram_freq=1  
    ) ]
```

Demonstration

Tensorboard

Generative Learning

???

???